


Gemma 4 E2B — KV-Cache Block Distribution by Layer



OrthoCache targets the 3 global attention layers (4, 9, 14) where the KV-cache grows with sequence length. The 12 sliding-window layers have fixed 512-token caches — only 1 block each, making eviction unnecessary.