

OrthoCache: Hardware-Native Multi-Band Spectral Attention Block Eviction on TPUs

Justin Arndt

Independent Researcher

justinarndt05@gmail.com

<https://github.com/j-arndt/orthocache>

Abstract—We present OrthoCache, an inline KV-cache eviction governor for Transformer attention on TPU accelerators. OrthoCache uses a Fast Walsh–Hadamard Transform (FWHT) to decompose key blocks into discrete sequency bands, enabling both a query-aware attention logit bound and a query-independent *Spectral Decay Ratio* (ζ) that distinguishes semantically coherent blocks from noise-dominated blocks with identical spatial variance. We directly address the Parseval no-op objection: aggregate block energy is computable in $O(b)$ without the transform, but per-band energy decomposition is not—making the $O(b \log b)$ FWHT genuinely load-bearing. We prove a formal Total Variation (TV) distance bound demonstrating that the attention approximation error decays exponentially in the gap between the eviction threshold and the maximum retained logit. The bound is machine-checked in Lean 4 with zero `sorry` stubs. We implement OrthoCache as JAX/Pallas kernels with online softmax accumulation and evaluate on Gemma 4 31B running on a TPU v5e-8 pod. Across all 60 model layers, we observe zero truncation bound violations, FWHT Parseval identity to $< 3 \times 10^{-7}$ relative error, and a ζ separability ratio of 4.04×10^{12} between blocks with identical spatial variance—confirming the spectral transform is non-redundant. The current Pallas kernel achieves full memory savings (50% KV-cache reduction) but does not yet translate to wall-clock latency improvement due to the MXU’s inability to skip masked tile operations.

Index Terms—KV-cache optimization, attention sparsity, Walsh–Hadamard transform, TPU, Pallas kernels, online softmax, formal verification, spectral filtering

I. INTRODUCTION

Long-context Transformer inference is heavily memory-bandwidth-bound during the autoregressive decoding phase. As context windows scale beyond 128K tokens, streaming the Key-Value (KV) cache from High Bandwidth Memory (HBM) to on-chip Vector Memory (VMEM) becomes the primary latency and energy bottleneck, eclipsing the cost of the attention computation itself.

Existing KV-cache optimization methods fall into two paradigms. **Passive compression**—quantization (e.g., TurboQuant [1]) and low-rank projection—reduces the storage footprint but retains all tokens in the pipeline. **Token-level eviction**—Heavy-Hitter architectures such as H₂O [2] and StreamingLLM [3]—drops tokens but requires query-key attention scores to determine relevance, introducing a circular dependency that prevents pre-dispatch eviction.

Furthermore, naive block-eviction schemes based only on key norm/energy run into a fundamental mathematical constraint: by

Parseval’s identity, applying an orthogonal transform preserves the Frobenius norm:

$$\|\hat{K}_j\|_F^2 = \|K_{B_j}\|_F^2 \quad (1)$$

Thus, computing aggregate block energy via the FWHT is a *Parseval no-op*: the transform adds $O(b \log b)$ computation for a quantity available in $O(b)$ via the bias-variance identity $\sum_{i \in B_j} \|k_i - \mu_j\|_2^2 = \sum_{i \in B_j} \|k_i\|_2^2 - b \|\mu_j\|_2^2$.

OrthoCache subverts this no-op through **Multi-Band Sequency Filtering**. Rather than collapsing the Walsh–Hadamard spectrum into a single aggregate energy scalar, we partition the $b = 512$ spectral coefficients into discrete frequency bands and compute the *Spectral Decay Ratio* ζ —the ratio of high-frequency to low-frequency energy. This per-band decomposition is genuinely uncomputable from spatial statistics: two blocks with identical total variance can have radically different ζ values. Combined with a query-aware logit upper bound, OrthoCache uses a **two-gate eviction criterion**: blocks must pass both a query-relevance gate and a spectral coherence gate.

A. Contributions

- 1) A **query-aware spectral bound** that uses the DC component for query alignment and the aggregate AC energy for residual bounding (Section II-A).
- 2) **Multi-Band Sequency Filtering** with the Spectral Decay Ratio $\zeta = E_{\text{high}}/E_{\text{low}}$ —a query-independent entropy signature that distinguishes semantically coherent blocks from noise-dominated blocks, making the FWHT genuinely load-bearing (Section II-B).
- 3) A **compilable TPU kernel** in JAX/Pallas using a single-axis grid (`num_heads,`) and online softmax accumulation, with pre-matmul mask gating that zeros evicted k/v blocks before the dot product (Section II-D).
- 4) **Formal proofs** of the Parseval identity and the exponential TV distance bound, type-checked in Lean 4 with zero `sorry` stubs (Section III).

II. METHOD

A. Query-Aware Spectral Bounds

Let $K \in \mathbb{R}^{N \times d_k}$ be the keys in the cache and $q \in \mathbb{R}^{d_k}$ be the current query vector. We partition K along the sequence axis into $m = N/b$ blocks $B_1, \dots, B_m$ of size b ($b = 512$

TABLE I
SEQUENCY BAND PARTITIONING FOR $b = 512$ WALSH COEFFICIENTS.

Band	Indices	Count	Interpretation
DC	0	1	Block mean (macro-semantic pivot)
Low-seq.	1–63	63	Smooth semantic trends
Mid-seq.	64–255	192	Syntactic/relational context
High-seq.	256–511	256	Rapid oscillations / noise

throughout). For each block B_j , the Fast Walsh–Hadamard Transform yields:

$$\hat{K}_j = \frac{1}{\sqrt{b}} \mathcal{H}_b K_{B_j} \quad (2)$$

where $\mathcal{H}_b$ is the $b \times b$ Hadamard matrix. We decompose $\hat{K}_j$ into its 0th coefficient (the *DC component* $\hat{K}_{j,0}$) and the remaining *AC components* $\hat{K}_{j,1:b-1}$.

The DC component represents the scaled block mean:

$$\mu_j = \frac{1}{b} \sum_{i \in B_j} k_i = \frac{1}{\sqrt{b}} \hat{K}_{j,0} \quad (3)$$

The AC components capture the variance from the mean. By Parseval’s identity:

$$\sum_{i \in B_j} \|k_i - \mu_j\|_2^2 = \sum_{s > 0} \|\hat{K}_{j,s}\|_2^2 \triangleq E_{j,AC} \quad (4)$$

For any token $i \in B_j$, the raw attention logit $z_i = q^\top k_i / \sqrt{d_k}$ can be bounded via Cauchy–Schwarz:

$$z_i \leq \underbrace{\frac{q^\top \hat{K}_{j,0}}{\sqrt{b \cdot d_k}}}_{\text{query-mean alignment}} + \underbrace{\frac{\|q\|_2}{\sqrt{d_k}} \sqrt{\frac{E_{j,AC}}{b}}}_{\text{residual bound}} \triangleq \tau_j(q) \quad (5)$$

Block B_j is a *candidate for eviction* if $\tau_j(q) < \tau$ for threshold τ . The DC term provides query alignment; the AC term bounds the residual—making the WHT essential for the query-aware bound.

B. Multi-Band Sequency Filtering

The aggregate AC energy $E_{j,AC}$ collapses the entire non-DC spectrum into a single scalar. Critically, this scalar equals the spatial variance $\sum_{i \in B_j} \|k_i - \mu_j\|_2^2$ and is therefore computable in $O(bd_k)$ without the $O(bd_k \log b)$ FWHT—making the transform redundant for this quantity alone.

To make the FWHT *genuinely load-bearing*, we partition the $b = 512$ Walsh coefficients into discrete **sequency bands** (Table I).

In the Walsh–Hadamard basis, *sequency*—the number of sign changes along a basis vector—plays the role of frequency. Low-sequency basis vectors produce smooth step-function patterns that capture coherent semantic dependencies; high-sequency vectors oscillate rapidly, capturing formatting noise and punctuation patterns.

We define the per-band energies:

$$E_j^{\text{low}} = \sum_{s=1}^{63} \|\hat{K}_{j,s}\|_2^2, \quad E_j^{\text{high}} = \sum_{s=256}^{511} \|\hat{K}_{j,s}\|_2^2 \quad (6)$$

and the **Spectral Decay Ratio**:

$$\zeta_j = \frac{E_j^{\text{high}}}{E_j^{\text{low}} + \epsilon_{\text{stab}}} \quad (7)$$

where $\epsilon_{\text{stab}} = 10^{-6}$ prevents division by zero.

Interpretation. High $\zeta_j (\gg 1)$ indicates that the block’s variance is dominated by high-frequency sign oscillations—characteristic of formatting tokens, punctuation sequences, and structural noise. Low $\zeta_j (\ll 1)$ indicates energy concentrates in the structural low-frequency bands—characteristic of continuous human language and long-range semantic dependencies.

1) *Why ζ is Uncomputable from Spatial Statistics:* Two blocks can have identical total spatial variance $\sum_i \|k_i - \mu_j\|_2^2$ but entirely different spectral decay ratios. A block of repetitive JSON formatting tokens (`{`, `"`, `}`) and a block of coherent technical prose may share the same Frobenius norm, but their frequency decompositions are radically different.

Theorem 1 (Non-reducibility of ζ). *There exist blocks A and B with $E_A^{\text{AC}} = E_B^{\text{AC}}$ (identical spatial variance) but $\zeta_A \neq \zeta_B$. Therefore, no spatial-domain function $f(\{k_i\}_{i \in B_j})$ that depends only on aggregate statistics (mean, variance, norms) can compute ζ_j without access to the individual spectral coefficients provided by the FWHT.*

Proof. By construction. Let block A be the weighted sum of Walsh basis vectors with indices 1–15 (all within the low-sequency band). Then $E_A^{\text{high}} = 0$ and $\zeta_A = 0$. Let block B be i.i.d. Gaussian noise scaled to have $E_B^{\text{AC}} = E_A^{\text{AC}}$. By the spectral flatness of white noise, energy distributes approximately uniformly across all 511 AC coefficients, so $\zeta_B \approx 256/63 \approx 4.06$. This is verified empirically in `tests/test_spectral_bands.py`. $\square$

2) *Query-Independence and Pipeline Hoisting:* ζ_j is *deliberately query-independent*. It is a structural property of the key block, computed once when the block is written to the KV cache and stored as a static scalar metadata entry. This enables:

- 1) **Asynchronous computation:** The FWHT and ζ computation execute during the initial KV-cache fill, amortized over thousands of subsequent decoding steps.
- 2) **PrefetchScalarGridSpec pipeline:** The pre-computed boolean mask streams into Scalar SRAM (SMEM) via `PrefetchScalarGridSpec` in parallel with query vector loading, introducing zero additional latency.
- 3) **No circular dependency:** Query-modulated ζ would force re-running the FWHT across the entire cache on every decoding step, reintroducing the $O(Nd_k \log b)$ tax.

Query-awareness is achieved *without* query-dependent transforms: the query norm $\|q\|_2$ dynamically scales the eviction threshold τ at runtime, tightening or loosening retention based on the current query’s activation magnitude.

3) *Two-Gate Eviction Criterion:* OrthoCache retains block B_j if and only if it passes **both** gates:

- 1) **Query-aware logit bound:** $\max_q \tau_j(q) \geq \tau$

2) Spectral coherence: $\zeta_j \leq \zeta_{\max}$

Gate 1 ensures blocks that could produce large attention logits for some query are retained. Gate 2 ensures blocks whose variance is dominated by high-frequency noise are evicted regardless of total energy—the key innovation that spatial variance cannot replicate.

C. Truncation Bound

Theorem 2 (OrthoCache Truncation Bound). *Let S denote the set of retained token indices and S^c the evicted set. Let α be the full attention distribution and $\hat{\alpha}$ the truncated distribution re-normalized over S . If all evicted blocks satisfy $\tau_j(q) < \tau$, then:*

$$\text{TV}(\alpha, \hat{\alpha}) \leq |S^c| \cdot \exp(\tau - z_{\max}) \quad (8)$$

where $z_{\max} = \max_{j \in S} z_j$ is the maximum logit among retained tokens.

Proof. The Total Variation distance equals the total evicted softmax mass: $\text{TV}(\alpha, \hat{\alpha}) = \sum_{i \in S^c} \alpha_i \triangleq \delta$. Each evicted token $i \in S^c$ satisfies $z_i \leq \tau_j(q) < \tau$. Therefore:

$$\alpha_i = \frac{e^{z_i}}{Z} \leq \frac{e^\tau}{e^{z_{\max}}} = \exp(\tau - z_{\max}) \quad (9)$$

Summing over the $|S^c|$ evicted tokens yields the bound. $\square$

The bound is *exponentially tight*: for typical attention distributions where $z_{\max} - \tau \gg 1$ (i.e., the retained tokens have much higher logits than the eviction threshold), the TV distance decays exponentially.

D. TPU Kernel with Online Softmax

To compile on TPU v5e using the JAX/Pallas kernel language, we structure the sparse attention kernel as follows:

- 1) **Grid size (`num_heads`,)**: One thread group per attention head eliminates write-after-write race conditions present when gridding over both blocks and heads.
- 2) **Sequential block loop**: The kernel loops over all m blocks. For each block, the boolean mask value is loaded.
- 3) **Pre-matmul mask gating**: For masked (evicted) blocks, the key and value tensors are zeroed *before* the matrix multiply:

$$k_{\text{active}} = \begin{cases} k_{\text{block}} & \text{if mask}[b] = \text{True} \\ 0 & \text{otherwise} \end{cases} \quad (10)$$

This ensures zero data contribution from evicted blocks at the operand level. The MXU instruction still fires (TPU’s Matrix Unit does not branch), but the zero-operand multiply is data-correct.

- 4) **Online softmax accumulation**: Running accumulators ($r_{\max}$, r_{sum} , r_{out}) are maintained per query token. Updates use `jnp.where` for TPU-traceable conditional execution without host round-trips.

Limitation. True FLOP elision requires either `pl.when()` support in Pallas (on the roadmap) or an XLA HLO reindexing pass that removes masked blocks from the loop entirely.

TABLE II

ζ DISTRIBUTION ACROSS LAYER TYPES (GEMMA 4 31B, 60 LAYERS).

Layer Type	Count	ζ Mean	ζ Std	ζ Range
Global	10	5.45	0.40	[4.76, 6.01]
Sliding	50	5.65	1.07	[4.65, 10.44]

III. LEAN 4 FORMALIZATION

The mathematical framework is formally verified in Lean 4 using Mathlib v4.8.0:

- `ParsevalWHT.lean` defines the Walsh–Hadamard matrix orthogonality and proves the L^2 norm preservation (Parseval’s identity for $\mathcal{H}_b$).
- `TruncationBound.lean` proves the TV distance reduction and the exponential softmax bound (Theorem 2).

All proofs compile cleanly via `lake build` with **zero sorry stubs**, verifying the mathematical foundation end-to-end.

IV. EXPERIMENTAL RESULTS

We evaluate OrthoCache on a Kaggle TPU v5e-8 pod using Gemma 4 31B (30.7B parameters, 60 layers, 16 sliding KV heads + 4 global KV heads).

A. Correctness Validation

All JAX-compiled query-aware bounds and masks match our reference NumPy calculations within bfloat16 numerical tolerance. The FWHT Parseval identity is verified across three representative layers: sliding (L0, dim $256 \rightarrow 512$), global early (L5, dim 512), and global late (L59, dim 512), achieving relative error $< 3 \times 10^{-7}$ (Parseval) and $< 4 \times 10^{-8}$ (JAX vs. NumPy reference). Truncation bound violations: **0** across all runs.

The critical test `test_zeta_not_computable_spatially` constructs two blocks with identical spatial variance ($E_A^{\text{AC}} = E_B^{\text{AC}}$, block A variance = 30,581.07, block B variance = 30,581.07) but $\zeta_A = 0.0000$ and $\zeta_B = 4.0373$, with a separation ratio of 4.04×10^{12} —empirically confirming Theorem 1.

Across all 60 layers, global layers exhibit $\zeta_{\text{mean}} = 5.45 \pm 0.40$ while sliding layers show $\zeta_{\text{mean}} = 5.65 \pm 1.07$ (Table II). Deeper layers (L30–45) concentrate around $\zeta \approx 4.8$, while early layers (L0–12) range from 6.5 to 10.4.

The truncation bound (Theorem 2) holds with zero violations across all tested eviction rates (9.4%–68.8%), with reconstruction error bounded below 2% even at 50% eviction (Table IV).

B. Latency Analysis

The current kernel provides **correctness but not acceleration** (Table III). It processes all blocks and gates accumulation via `jnp.where`; the MXU executes the matmul regardless. True performance gains require hardware-level block skipping.

The $1.05\times$ speedup at 8K tokens (2.31 ms vs. 2.44 ms) is within measurement noise ($\sigma \approx 0.06$ – 0.12 ms). At all other sequence lengths, sparse and dense attention are within 0.1% of each other.

TABLE III

LATENCY COMPARISON: DENSE VS. SPARSE ATTENTION (50% EVICTION).

Seq. Len.	Blocks	Dense (ms)	Sparse (ms)	Speedup
4,096	8	2.381	2.393	0.99×
8,192	16	2.435	2.310	1.05×
16,384	32	2.999	3.000	1.00×
32,768	64	5.704	5.706	1.00×
65,536	128	11.005	11.017	1.00×

TABLE IV

TV DISTANCE VALIDATION (LAYER 5, GLOBAL ATTENTION, 50% TARGET EVICTION).

Target Evict.	Actual Evict.	TV Distance	Recon. Error
10%	9.4%	0.094	0.0026
30%	28.1%	0.281	0.0101
50%	50.0%	0.500	0.0184
70%	68.8%	0.688	0.0171

C. Cost-Benefit Projections

We include a parameterized infrastructure model that translates block sparsity into projected annual fleet-level savings (Table V). These figures assume the FLOP-skip kernel is deployed.

Caveat. The current prototype kernel achieves latency parity with dense attention (speedup $\approx 1.00\times$ across all tested sequence lengths). While this confirms zero overhead from the masking approach, the economic value materializes only when the sparse kernel achieves net-positive throughput via hardware-level block skipping (`pl.when()` or XLA HLO reindexing).

V. RELATED WORK

KV-Cache Compression. TurboQuant [1] applies mixed-precision quantization to the KV-cache, achieving 2–4× memory reduction. KIVI [4] introduces per-channel key quantization with per-token value quantization. These methods are orthogonal to OrthoCache: quantization reduces *per-token* storage while our approach reduces the *number of tokens* in the pipeline.

Token-Level Eviction. H₂O [2] identifies “Heavy Hitter” tokens via cumulative attention scores. StreamingLLM [3] discovers “attention sinks” in initial tokens. Both require attention score computation before eviction, creating a circular dependency. OrthoCache evicts before attention is computed by using spectral bounds.

Structured Sparsity. BigBird [5] and Longformer [6] impose fixed sparse attention patterns (local + global windows). These are static and content-independent. OrthoCache’s spectral filtering is content-adaptive: the eviction set depends on the actual information content of each block.

Walsh–Hadamard in ML. QuaRot [7] uses randomized Hadamard rotations to flatten activation outliers before quantization. QuIP# [8] applies incoherence processing via Hadamard matrices for weight quantization. Both operate *per-coordinate* on individual activations; OrthoCache operates *per-block* on key

TABLE V

PROJECTED ANNUAL FLEET-LEVEL SAVINGS († = PROJECTED).

Scenario	Sparsity	OpEx	CapEx	Total
Conservative	0.25	\$2.8M	\$20M	\$22.8M
Moderate	0.50	\$5.6M	\$60M	\$65.6M
Aggressive	0.70	\$7.8M	\$100M	\$107.8M

sequences and extracts per-band energy rather than performing coordinate-level rotation.

VI. CONCLUSION

OrthoCache resolves the Parseval no-op through Multi-Band Sequency Filtering. Rather than collapsing the Walsh–Hadamard spectrum into a redundant aggregate energy scalar, OrthoCache extracts per-band energy decompositions that distinguish semantically coherent blocks from noise-dominated blocks—a distinction impossible to make with spatial statistics alone. The Spectral Decay Ratio ζ provides a query-independent entropy signature computed once per block, while the query-aware logit bound provides dynamic, per-query relevance filtering. Together, these form a two-gate eviction criterion backed by a formally verified exponential TV distance bound.

A. Limitations

The current Pallas kernel provides correctness but not acceleration—it computes attention for all blocks and gates accumulation rather than skipping FLOPs. True hardware-level block skipping requires either `pl.when()` support in Pallas or an XLA HLO reindexing pass. The economic value projections assume this capability is realized.

B. Future Work: HLO Loop-Reindexing Pass

The primary engineering target is a native XLA `HloModulePass` that transforms the attention loop from *predication* (executing all blocks and suppressing write-back) to *compaction* (iterating only over retained blocks). The pass has three stages:

- 1) **Stream compaction pre-pass:** A parallel prefix sum on the boolean block mask produces an indirection lookup table of length K (the number of active blocks, where $K = (1 - S) \cdot M$). On the TPU VPU, prefix sums over 512-element boolean masks complete in $O(\log b)$ cycles.
- 2) **Dynamic coordinate inversion:** The pass rewrites the attention loop boundary from $[0, M)$ to $[0, K)$ and inserts indirect DMA fetches: `block_idx = active_indices[i]`. The MXU systolic array never touches evicted blocks.
- 3) **Irregular AllToAll collectives:** In multi-host tensor-parallel inference, the pass replaces the standard AllToAll with AllToAllv (variable-size shards), transmitting only the $(1 - S)$ fraction of active blocks across the ICI fabric—physically preventing evicted data from saturating bisection bandwidth.

Additional future directions include:

- **Adaptive $\zeta_{\max}$ calibration:** Learning $\zeta_{\max}$ per layer and per head from a calibration set. Our measurements show global layers have tighter ζ distributions ($\sigma = 0.40$) than sliding layers ($\sigma = 1.07$), suggesting per-layer-type thresholds.
- **Composable quantization:** Combining OrthoCache’s block eviction with per-token quantization (e.g., TurboQuant) for multiplicative memory savings.
- **Cross-model evaluation:** Extending validation to Llama 3, Mixtral, and architectures with different GQA configurations.

REFERENCES

- [1] Google DeepMind, “TurboQuant: Quantized KV-cache for efficient inference,” Internal report, 2025.
- [2] Z. Zhang, Y. Sheng, T. Zhou, T. Chen, L. Zheng, R. Cai, Z. Song, Y. Tian, C. Ré, C. Barrett, Z. Wang, and B. Chen, “H₂O: Heavy-Hitter Oracle for efficient generative inference of large language models,” in *Proc. NeurIPS*, 2023.
- [3] G. Xiao, Y. Tian, B. Chen, S. Han, and M. Lewis, “Efficient streaming language models with attention sinks,” in *Proc. ICLR*, 2024.
- [4] Z. Liu, J. Yuan, H. Jin, S. Zhong, Z. Xu, V. Braverman, B. Chen, and X. Hu, “KIVI: A tuning-free asymmetric 2bit quantization for KV cache,” in *Proc. ICML*, 2024.
- [5] M. Zaheer, G. Guruganesh, K. A. Dubey, J. Ainslie, C. Alberti, S. Ontanon, P. Pham, A. Ravula, Q. Wang, L. Yang, and A. Ahmed, “Big Bird: Transformers for longer sequences,” in *Proc. NeurIPS*, 2020.
- [6] I. Beltagy, M. E. Peters, and A. Cohan, “Longformer: The long-document transformer,” *arXiv preprint arXiv:2004.05150*, 2020.
- [7] S. Ashkboos, A. Mohtashami, M. L. Croci, B. Li, M. Jaggi, D. Alistarh, T. Hoeftler, and J. Hensman, “QuaRot: Outlier-free 4-bit inference in rotated LLMs,” *arXiv preprint arXiv:2404.00456*, 2024.
- [8] A. Tseng, J. Chee, Q. Sun, V. Kuleshov, and C. De Sa, “QuIP#: Even better LLM quantization with Hadamard incoherence and lattice codebooks,” in *Proc. ICML*, 2024.